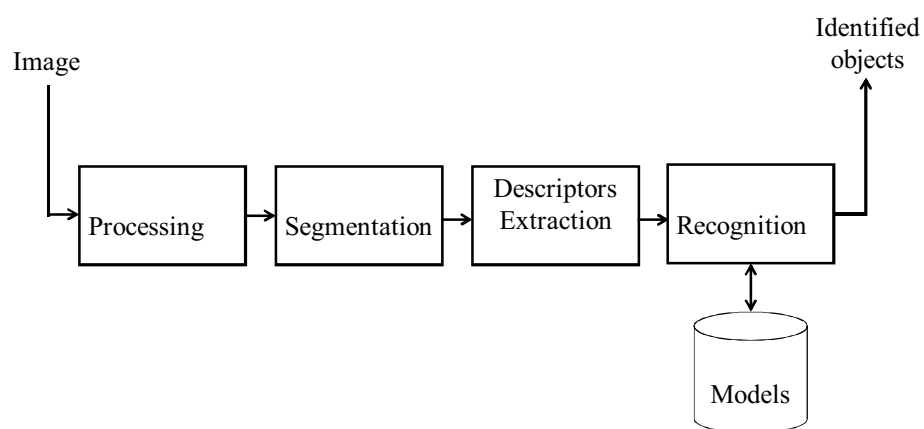


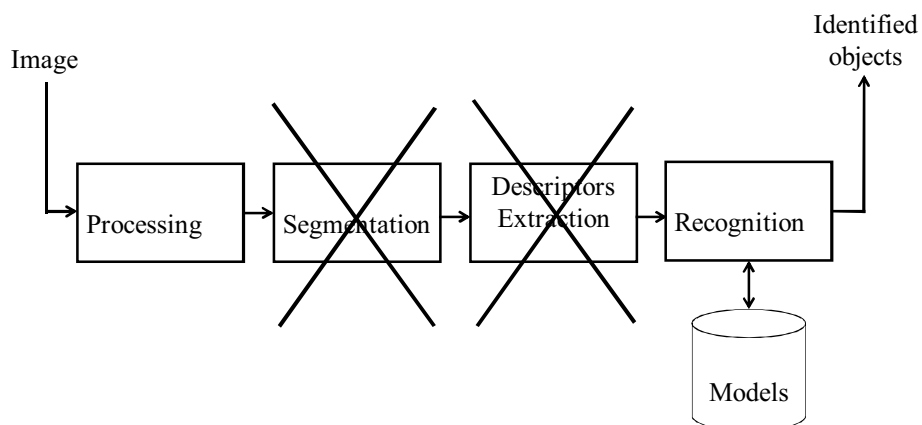
## Local Features



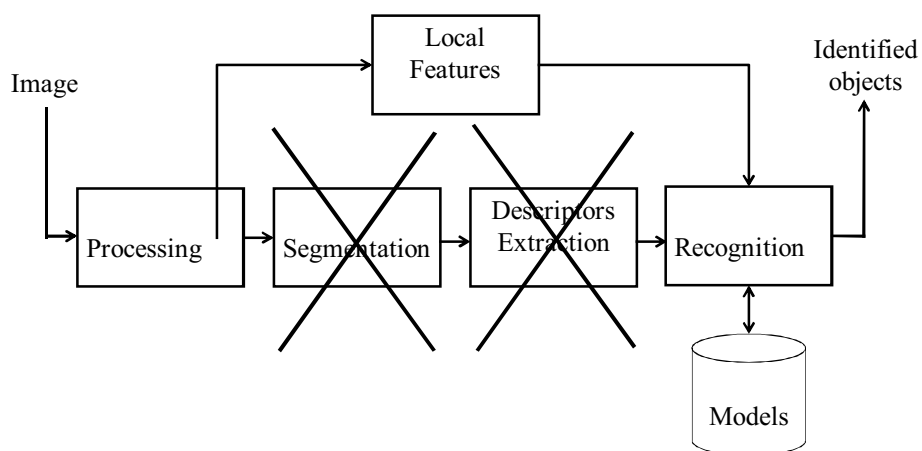
## A computer vision system



## A computer vision system



## A computer vision system



## Local Features

- Histograms
- Hough transform
- Corners
- Scale Invariant Feature transform (SIFT)
- Haar Features (face detection)



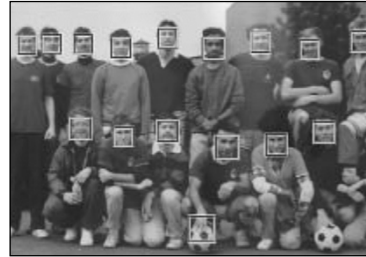
## Haar features

- Real-time X detection
- Where X can be:
  - Faces
  - Traffic signs
  - Cats
  - Any object



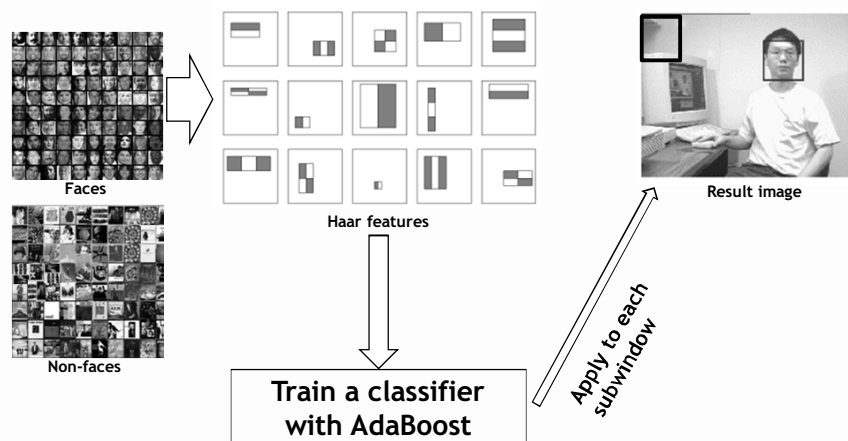
## Robust Real-time Object Detection

Paul Viola and Michael Jones



High-speed face detection with good accuracy

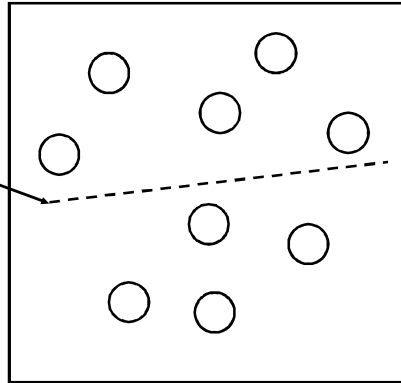
## Viola-Jones Face Detector



- Train with 5K positives, 350M negatives
- [Implementation available in OpenCV:  
<http://www.intel.com/technology/computing/opencv/>]

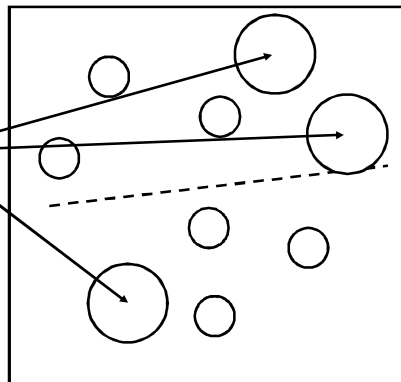
## The Adaboost classifier

Weak  
Classifier 1

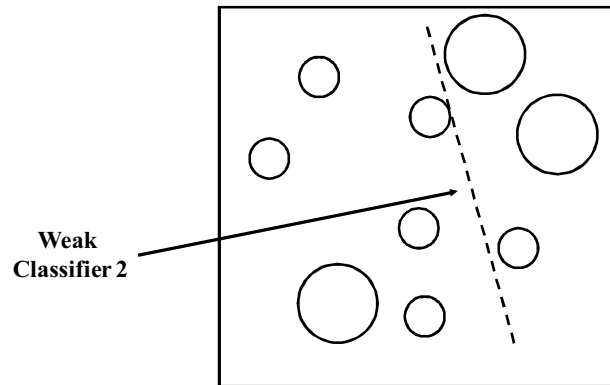


## The Adaboost classifier

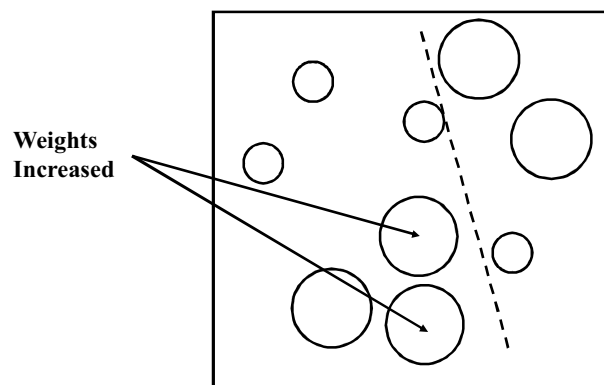
Weights  
Increased



## The Adaboost classifier

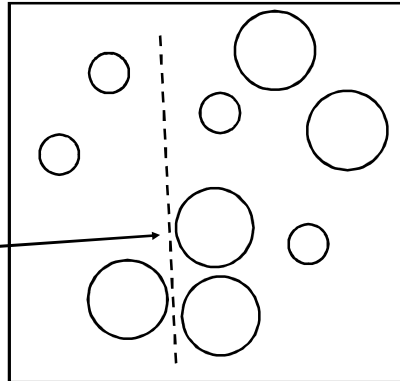


## The Adaboost classifier



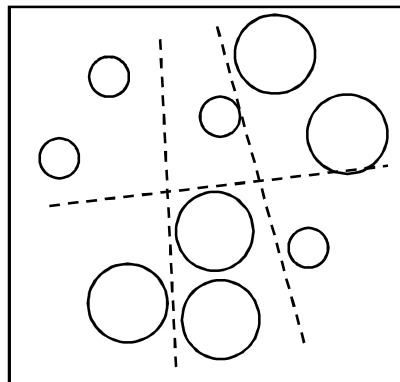
## The Adaboost classifier

Weak  
Classifier 3



## The Adaboost classifier

Final classifier is  
a combination of weak  
classifiers



## Boosting Algorithm

- Given example images  $(x_1, y_1), \dots, (x_n, y_n)$  where  $y_i = 0, 1$  for negative and positive examples respectively.
- Initialize weights  $w_{1,i} = \frac{1}{2m}, \frac{1}{2l}$  for  $y_i = 0, 1$  respectively, where  $m$  and  $l$  are the number of negatives and positives respectively.
- For  $t = 1, \dots, T$ :

1. Normalize the weights,

$$w_{t,i} \leftarrow \frac{w_{t,i}}{\sum_{j=1}^n w_{t,j}}$$

so that  $w_t$  is a probability distribution.

2. For each feature,  $j$ , train a classifier  $h_j$  which is restricted to using a single feature. The error is evaluated with respect to  $w_t$ ,  $\epsilon_j = \sum_i w_{t,i} |h_j(x_i) - y_i|$ .
3. Choose the classifier,  $h_t$ , with the lowest error  $\epsilon_t$ .
4. Update the weights:

$$w_{t+1,i} = w_{t,i} \beta_t^{1-e_i}$$

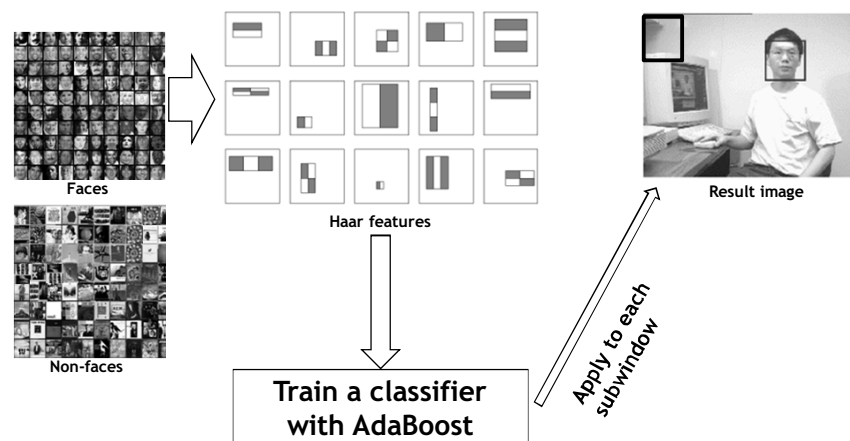
where  $e_i = 0$  if example  $x_i$  is classified correctly,  $e_i = 1$  otherwise, and  $\beta_t = \frac{\epsilon_t}{1-\epsilon_t}$ .

- The final strong classifier is:

$$h(x) = \begin{cases} 1 & \sum_{t=1}^T \alpha_t h_t(x) \geq \frac{1}{2} \sum_{t=1}^T \alpha_t \\ 0 & \text{otherwise} \end{cases}$$

where  $\alpha_t = \log \frac{1}{\beta_t}$

## Viola-Jones Face Detector



- Train with 5K positives, 350M negatives
- [Implementation available in OpenCV:  
<http://www.intel.com/technology/computing/opencv/>]

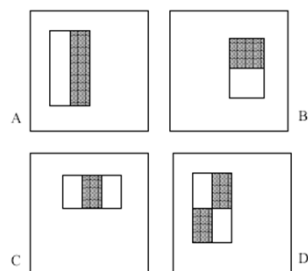


## Face Detection System

- Training Data
  - 5000 faces
    - All frontal, rescaled to 24x24 pixels
  - 9500 million non-faces
  - Faces are normalized
    - Scale, translation
- Many variations
  - Across individuals
  - Illumination
  - Pose



## Haar Features



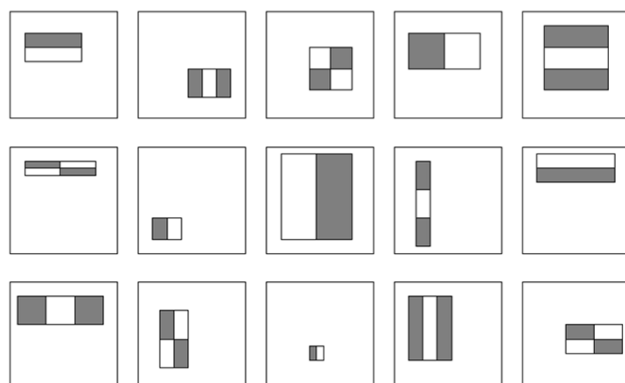
Rectangular filters



$$f(x, y) = \sum_i p_b(i) - \sum_i p_w(i)$$

Local features: Subtract sum of pixels in white area from the sum of pixels in black area;

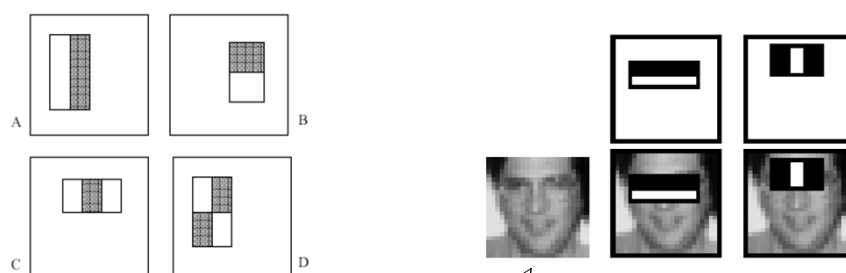
## Large library of filters



Considering all possible filter parameters:  
position, scale,  
and type:

180,000+  
possible features  
associated with  
each 24 x 24  
window

## Image Features



Rectangular filters

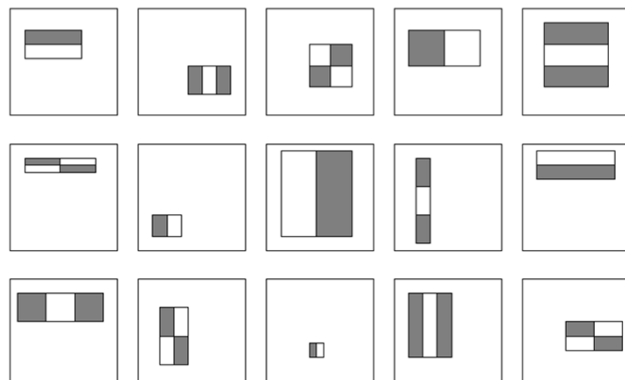
Local features: Subt

Too many features  
In a 24 x 24 patch with 4x4  
detector, there are over 160,000  
locations for rectangles

$$f(x) = \sum_i p_b(i) - \sum_i p_w(i)$$

of pixels in black area

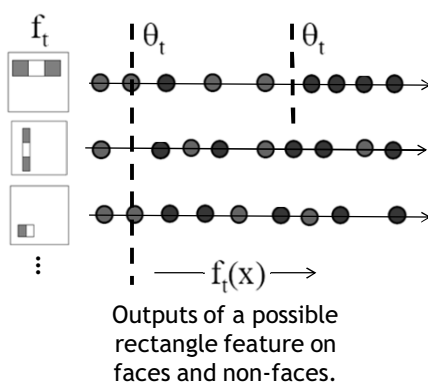
## Large library of filters



Considering all possible filter parameters:  
position, scale, and type:  
180,000+ possible features associated with each 24 x 24 window

## AdaBoost for feature+classifier selection

- Want to select the single rectangle feature and threshold that best separates positive (faces) and (non-faces) training examples, in terms of *weighted* error.



Resulting weak classifier:

$$h_t(x) = \begin{cases} +1 & \text{if } f_t(x) > \theta_t \\ -1 & \text{otherwise} \end{cases}$$

For next round, reweight the examples according to errors, choose another filter/threshold combo.

## The detector

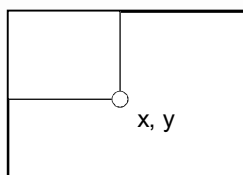
- A simple filter bank with learned weights applied across the image
- But with some notable performance-boosting implementation tricks...

## Three big speed gains

- Integral image representation and rectangle features
- Selection of a small but effective feature set with AdaBoost
- Cascading simple detectors to quickly eliminate false positives

## The integral image representation

An image representation that stores the sum of the intensity values above and to the left of the image point.

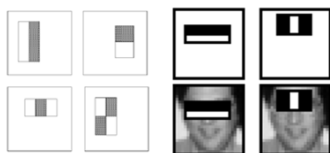


$\text{IntegralImage}(x,y) = \text{Sum of the values in the grey region}$

So what's it good for?

## The integral image

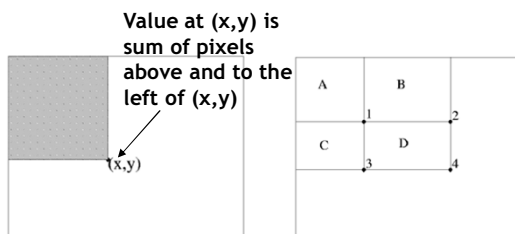
**“Rectangular” filters**



Feature output is difference between adjacent regions

**Efficiently computable with integral image: any sum can be computed in constant time**

**Avoid scaling images → scale features directly for same cost**



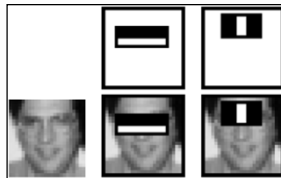
$$\begin{aligned}
 D &= 1 + 4 - (2 + 3) \\
 &= A + (A + B + C + D) - (A + C + A + B) \\
 &= D
 \end{aligned}$$

## Speed gain number two: AdaBoost selected features

AdaBoost is used to select the best set of rectangular features.

AdaBoost iteratively trains a classifier by emphasizing misclassified training data.

Assigned feature weights are used to select the “most important” features.



Top two features weighted by AdaBoost

## Intermediate results

The face detector using 200 AdaBoost-selected features achieved a 1 in 14084 false positive rate when turned for a 95% classification rate.

An 384x288 image took 0.7 seconds to scan.

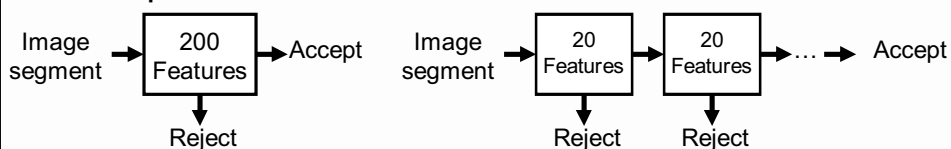
There are more improvements to be made...

## Speed gain number three: Cascading detectors

Instead of applying all 200 filters at every location in the image, train several simpler classifiers to quickly eliminate easy negatives.

Each successive filter can be trained on true positives and the false positives passed by the filters before it.

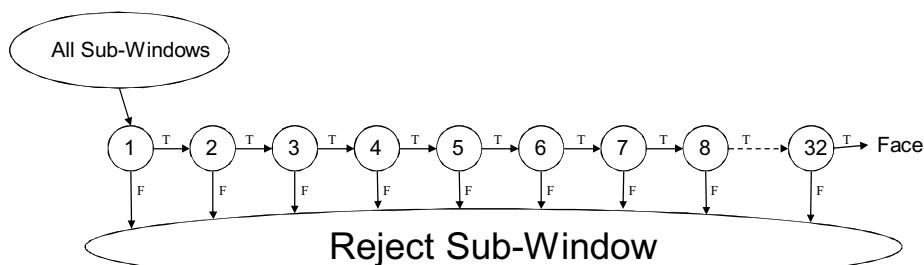
The filters are trained to allow approximately 10% false positives.



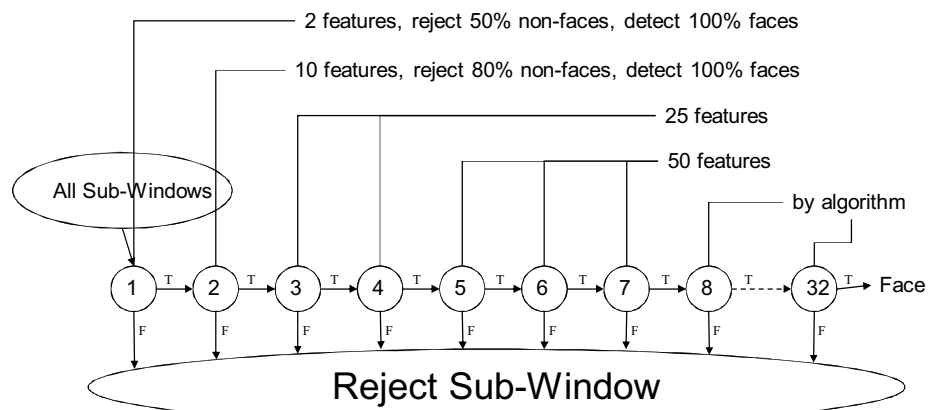
## Structure of the Detector Cascade

Combining successively more complex classifiers in cascade

- 32 stages
- included a total of 4297 features



## Structure of the Detector Cascade



## Viola-Jones Face Detector: Results





## Viola-Jones Face Detector: Results

